

- Goal of Pruning: Remove neurons that aren't as involved in producing results
 - Simple solution: Remove the connections of neurons that aren't being activated. This results in sparse matrix multiplications.
- Miscellaneous Links
 - Song Han Guest Lecture at Stanford: <https://www.youtube.com/watch?v=eZdOkDtYMoo>
 - <https://songhan.mit.edu/>
 - Tensorflow Pruning: https://www.tensorflow.org/model_optimization/guide/pruning/pruning_with_keras
 - L1 Ridge Regression
 - <https://www.youtube.com/watch?v=Q81RR3yKn30>
 - Reduces variance of the model (sensitivity to input values)
 - Does so by adding a term to the loss function
 - The term is $\lambda \cdot (\text{sum of square of each weight})$
 - The effect is that weights tend to 0, meaning that there is a reduced variance of prediction with respect to input values
 - L2 Lasso Regression
 - <https://www.youtube.com/watch?v=NGf0voTMIcs>
 - Similar applications as ridge regression
 - Does so by adding a term to the loss function, just like ridge regression
 - The term is $\lambda \cdot (\text{sum of absolute value of weights})$
 - Difference is that L2 reduces weights closer to 0 by an equal magnitude, while L1 reduces weights asymptotically close to 0
 - Useful in high-dimensional datasets where only a fraction of the parameters are useful
- Ideas
 - RL for predicting which connections should be pruned, which should be kept
 - Dynamically removing and adding nodes to the computational graph
 - How does parameter initialization influence convergence and training?
 - Can we test out different initialization approaches and see which works?
 - Different methods of pruning (e.g. probabilistic pruning)
 - Paper *Learning both Weights and Connections for Efficient Neural Networks*, for example, tested different pruning types
 - "We also experimented with probabilistically pruning parameters based on their absolute value, but this gave worse results."
 - Random thought- can a conv layer be equivalently represented by series of two dense layers?
 - Having pruning work at intervals that decay over time
 - Ex: Prune first every 1 batch, then every 2 batches then every 4 batches
 - Playing around with pruning schedule
 - Regularizer function that encourages parameters to go to integer values rather than just 0
 - Dataset-dependent model initialization
 - Dude you have no idea turns out it's an actual thing

- <https://arxiv.org/pdf/1511.06856.pdf>
 - I wrote this idea down before hearing about it
 - Smart data augmentation
- Mark Kurtz, Pruning Deep Learning Models for Accuracy
 - Most neurons in neural networks are useless
 - Sparsity of model can be reduced 70-90%
 - Effect: Decreasing compute time, improving performance, decreasing memory
 - Structured vs unstructured
 - Structured - removing groups of weights & neurons
 - Unstructured - removing connections at random
 - Pruning has several hyperparameters
 - How often to prune
 - Which neurons to prune
 - Speaker presents automated solution for pruning that his team at NeuralMagic as developed
 - User interface to manipulate sparsity
 - API for Pytorch training that prunes every epoch
- Han, Learning both Weights and Connections for Efficient Neural Networks
 - <https://arxiv.org/pdf/1506.02626.pdf>
 - Takeaway: Baseline here is that this paper proposed the process of iterative training and pruning. Consequently, the model necessitates less compression, achieves smaller size, with little sacrifice in accuracy
 - How did they do it?
 - Essence of the paper: 3-step flow
 - Training, Pruning, Adjusting weights
 - Training provides insights into which weights to remove
- Han, Deep Compression
 - <https://arxiv.org/pdf/1510.00149.pdf>
 - <https://www.youtube.com/watch?v=CrDRr2fxbsg&t=2s>
 - Takeaway: Seems to be the pioneering paper for pruning of neural networks from back in 2015. Emphasis on model size reduction, as well as speed and energy consumption.
 - How did they do it?
 - 3 step procedure
 - Pruning weights with low connection magnitude
 - Make cluster of floating point weights with similar values
 - Approximate cluster of floating point weights with integer centroids
 - Fine tune centroids
 - Quantization of weights → weight sharing
 - Weight sharing implies a discrete set of weights → make lookup table
 - k clusters of weights (that share the same value) → $\log_2 k$ bits to encode index of each weight

- Given n connections, all connections can be encoded in $\log_2 k * n$ connections
 - Given b bits to represent each connection, lookup table is kb bits
 - Without lookup table, the neural network requires $n*b$ bits
- [Huffman encoding](#)
 - Efficient form of encoding, where key (encoded symbol) length depends on frequency of value in message
 - I guess this is an optional addition to step 2, which further increases the compression rate of the network
- Learning the number of neurons in neural networks
 - <https://arxiv.org/pdf/1611.06321.pdf>
 - https://www.cs.ubc.ca/~schmidtm/MLRG/structured_sparsity.pdf
 - <https://towardsdatascience.com/sparse-group-lasso-in-python-255e379ab892>
 - <https://www.youtube.com/watch?v=vgeDzliXhWc>
 - This pretty much explains it all
 - Takeaway: This paper proposes a method to cut down the number of neurons during training by means of sparsity regularization without the need for a preprocessing step
 - Questions
 - The authors mention that the groups are non-overlapping. Is this necessarily true, or is this just true in this case?
 - How did they do it?
 - They focus on learning the number of weights in a neural network
 - These guys just added a term to the loss function and are writing a paper
 - $L(\Theta) = L_0(\Theta) + r(\Theta)$
 - $r(\Theta)$ is a regularization term expressed as
 - $(\sum \lambda \sqrt{P_i})$ for every layer i * $(\sum |\theta^n_i|^2 \text{ for every input/output sample } N)$
 - P_i is a vector of parameters for every layer grouped together
 - Basically it just seems that they found a new regularization term that builds off the ideas of L1 and L2 regularization (notes at top of document)
 - L1 Ridge Regularization
 - $r = \lambda * (\text{sum of squares of weights})$
 - Benefit: Leads to small parameter values
 - L2 Lasso Regularization
 - $r = \lambda * (\text{sum of absolute value of weights})$
 - Benefit: Leads to sparse solutions (zeroed out weights)
 - Group Lasso
 - $r = (\text{sum of } (\lambda * \sqrt{p_i}) * \text{sum of } (\text{every parameter in the group})^2)$ for every non-overlapping group p_i
 - Sparse Group Lasso
 - Basically group lasso regularization term + L2 regularization term
 - Adding just one more term to group lasso

- Regularization in neural networks: A Taxonomy
 - <https://arxiv.org/pdf/1710.10686.pdf>
 - Reviewing and summarizing various types of regularization in neural networks
 - Regularization: A method that attempts to increase test dataset performance of a model
 - Types
 - Data
 - Adding noise to the parameters (e.g. Dropout)
 - Data augmentation
 - Changing representation of the input
 - Error function
- Learning to Prune Filters in Convolutional Neural Networks
 - <https://arxiv.org/pdf/1801.07365.pdf>
 - <https://www.youtube.com/watch?v=3yOZxmlBG3Y>
 - Takeaway: These guys found a way to prune the majority of filters with 0 loss in accuracy
 - Pruned 92% (!) of filters with just 3% loss in accuracy
 - How did they do it?
 - They created a separate neural network (NN_{prune}) that predicted which filters to remove.
 - NN_{prune} did not remove individual connections, as they said this did not boost performance by much.
 - NN_{prune} was trained using a reinforcement learning approach where the weights were modified via a policy gradient method.
 - NN_{prune} could not be trained using gradient descent because the reward function was non-differentiable
 - Other notes
 - The paper provides nice references to related work on model pruning techniques
- He, AMC: AutoML for Model Compression
 - <https://arxiv.org/pdf/1802.03494.pdf>
- J. Frankle, The Lottery Ticket Hypothesis
 - <https://arxiv.org/pdf/1803.03635.pdf>
 - <https://www.youtube.com/watch?v=s7DqRZVvRiQ>
 - Takeaway: The researchers investigate whether they can train sparse pruned neural networks from scratch.
 - The typical end-product of the cycle of pruning (removing weights with new magnitude, retraining) is a sparse neural network
 - Training this sparse neural network from scratch results in a neural network with lower accuracy
 - How did they do it?
 - These authors decided to train a dense neural network partially, then prune it

- They realize that retraining the sparse neural network with **the original initializations** results in a network that can train well
 - Retraining the pruned neural network from scratch (with same initializations) enables faster training, higher test accuracy
 - They call this sub-network the “lottery ticket”
 - But it’s also important to realize that initialization isn’t everything. Rearranging which connections were pruned does not result in final accuracy close to the original
- Han, Once for All
 - <https://arxiv.org/pdf/1908.09791.pdf>
 - https://www.youtube.com/watch?v=_jKImMePY-w